

Assignment 2 Report

Prayash Kumar Sahu (22B1261), Aditya Singh Bhadoria (22B1247)

1 System Architecture

1.1 Components

- **DeepFace (Facenet):** Generates face embeddings for recognition using the `Facenet` model.
- **OpenCV:** Handles camera capture, image processing (resizing, augmentation), frame manipulation, and real-time face detection (using its built-in detector).
- **NumPy:** Performs all numerical computations, including creating a trusted "centroid" embedding, calculating cosine similarity, and managing embedding arrays.
- **TensorFlow:** Serves as the underlying deep learning framework supporting DeepFace operations (v2.13.0).
- **Image Augmentation:** A pre-processing step creates robust embeddings by generating variations of source images, including horizontal flipping, brightness adjustments ($\times 0.8, \times 1.2$), rotations ($\pm 10^\circ$), and scaling ($\times 0.9, \times 1.1$).
- **Trusted Centroid:** The system does not store individual trusted embeddings for real-time comparison. Instead, it computes a single "average" embedding (centroid) from all augmented trusted faces.
- **Auto-Calibrated Threshold:** The classification boundary is not hard-coded. It is dynamically calculated as the midpoint between the average similarity of trusted faces to the centroid and the average similarity of random faces to the centroid.
- **Unknown Face Logging:** When an "UNKNOWN" face is detected, the system saves a cropped image of the face to the `unknown_faces/` directory with a timestamp and a 10-second cooldown period to prevent spam.

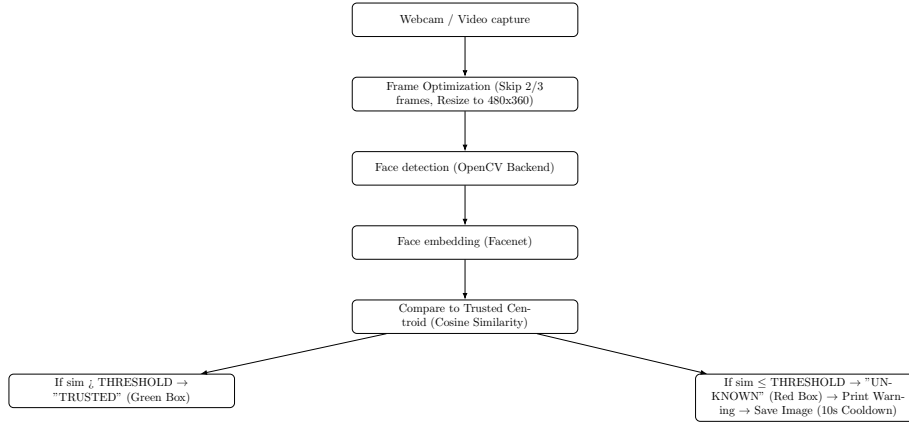


Figure 1: Flow diagram of the system architecture

2 Integration Challenges and Solutions

1. **Challenge:** Facial recognition is highly sensitive to variations in lighting, head pose, and distance, which can lead to false negatives for trusted users.

Solution: Implemented a robust pre-processing pipeline. Instead of a single embedding, the system generates multiple augmentations (rotation, scale, brightness, flip) for each trusted image, computes embeddings for all, and averages them into a single, highly robust vector. This "mean" embedding is much more resilient to real-world conditions.

2. **Challenge:** High-accuracy face detectors like **RetinaFace** are too computationally expensive for smooth, real-time webcam processing.

Solution: A hybrid detector strategy was used. The slow, high-accuracy **RetinaFace** detector is used **only** for the one-time, offline pre-processing of trusted and random faces. For the real-time application, the system switches to the much faster **opencv** backend for face detection.

3. **Challenge:** Real-time video processing (detection + embedding) can still cause significant lag, even with a faster detector.

Solution: Two performance optimizations were added to the main loop:

- **Frame Resizing:** The input frame is immediately resized to a small 480x360 resolution, drastically reducing the pixels the detector needs to scan.
- **Frame Skipping:** The system processes only one out of every three frames (`frame_count % 3`), effectively tripling the potential FPS by reducing the processing load.

4. **Challenge:** Manually selecting a fixed similarity threshold (e.g., 0.6) is unreliable and may not be optimal for the specific faces enrolled in the system.

Solution: The system uses an auto-calibrated threshold. It calculates the similarity of known trusted faces and known random faces (from the `random_faces` folder) against the trusted centroid, then sets the threshold as the midpoint between these two groups: $(\text{mean}(\text{trusted_sims}) + \text{mean}(\text{random_sims}))/2$.

3 Ethical Considerations and Testing Results

3.1 Ethical Considerations

The deployment of this system involves handling sensitive biometric data, raising key ethical points:

- **Privacy and Data Storage:** The system stores face embeddings (numerical representations) in `embeddings.npz`. More critically, it actively captures and saves images of "UNKNOWN" individuals to the `unknown_faces` directory. This requires a clear policy for data retention, access control, and deletion.
- **Informed Consent:** Individuals whose images are placed in the `trusted_faces` folder must provide explicit consent for their biometric data to be used for authentication.
- **Bias and Fairness:** The `Facenet` model, like all facial recognition models, may have performance biases across different demographics. The use of image augmentation (brightness, scaling) in the pre-processing step is a technique that can help mitigate some of these biases by creating a more generalized embedding.
- **Proportionality of Response:** The system's response to an "UNKNOWN" detection is relatively passive: a console warning ("You are not authorized!! Please Leave!") and logging the face. This is a less intrusive security measure than a loud alarm or immediate authority notification.

3.2 Testing Results

The notebook provides execution output from the pre-processing and initialization stages, which serve as the system's baseline configuration and calibration:

- **Database Configuration:** The pre-processing script (Cell 7) was executed with a database of 1 trusted face (which was augmented to generate multiple embeddings) and 5 random faces.

- **Auto-Calibrated Threshold:** Based on the 1 trusted and 5 random face embeddings, the real-time script (Cell 8) successfully auto-calibrated its classification boundary. The calculated cosine similarity threshold was **0.619**.
- **System Status:** The real-time detection script successfully initialized the webcam and began monitoring, confirming that all components (OpenCV, DeepFace, NumPy) were loaded correctly.
- **Performance Features:** The code demonstrates the implementation of robustness features (augmentation, mean embeddings) and real-time performance optimizations (frame skipping, resizing) in its design, though no quantitative latency or accuracy metrics (e.g., FPS, FRR/FAR) are provided in the notebook.

4 Instructions to Run Code

4.1 Prerequisites

- Python 3.8 or higher.
- A functional webcam connected to the system.
- Required Python packages: `deepface`, `opencv-python`, `speechrecognition`, `pyttsx3`, `pyaudio`, `tensorflow`.

4.2 Setup and Installation

1. Create and activate a Python virtual environment (recommended).
2. Install all required packages using pip:

```
pip install deepface opencv-python speechrecognition pyttsx3 pyaudio tensorflow
```
3. Create three folders in the same directory as the notebook: `trusted_faces/`, `random_faces/`, and `unknown_faces/`.
4. Place at least one high-quality image (JPG, PNG, WEBP) of a trusted person in the `trusted_faces/` folder.
5. Place several images of different, non-trusted people in the `random_faces/` folder. These are used to calculate the threshold.

4.3 Execution

The notebook is divided into two main parts that must be run in order.

1. **Generate Embeddings:** Run the code cell containing the `augment_image`, `get_embedding`, and `compute_all` functions (Cell 7). This script will process all images in `trusted_faces/` and `random_faces/` and save the results into a new file named `embeddings.npz`.
2. **Run Real-Time Verification:** After `embeddings.npz` has been created, run the next code cell (Cell 8). This script will:
 - Load the embeddings.
 - Calculate the auto-calibrated threshold (e.g., 0.619).
 - Open the webcam and display the "Face Verification" window.
3. **Operation:** The system will draw a green "TRUSTED" box around recognized faces and a red "UNKNOWN" box around others. Unknown faces will be saved to the `unknown_faces/` folder.
4. **To Stop:** Press the 'q' key while the "Face Verification" window is active.